

# **Data-Informed Global Sparseness in Attention Mechanisms for Deep Neural Networks**

<https://arxiv.org/abs/2012.02030>

Ileana Rugina\*, Rumen Dangovski\*, Li Jing, Preslav Nakov, Marin Soljačić

# Self-Attention

# Self-Attention Mechanisms



$$X \in \mathbb{R}^{N \times h}$$



$$Q = XW^Q, Q \in \mathbb{R}^{N \times d}$$

$$K = XW^K, K \in \mathbb{R}^{N \times d}$$



$$V = XW^V, V \in \mathbb{R}^{N \times d}$$

word  
embeddings

# Self-Attention Mechanisms

$$\text{softmax} \left( \frac{Q \times K^T}{\sqrt{d_k}} \right) V = Z$$

$$A = \text{softmax} \left( \frac{QK^T}{\sqrt{d}} \right), A \in \mathbb{R}^{N \times N}$$

$$Z = AV, Z \in \mathbb{R}^{N \times d}$$

# Related Work

# Kernel Trick Ideas

$$Z = \text{softmax} \left( \frac{QK^T}{\sqrt{d}} \right) V$$

$$\text{softmax} \left( \frac{QK^T}{\sqrt{d}} \right) = \phi(Q)\phi(K)^T$$

$$Z = (\phi(Q)\phi(K)^T) V = \phi(Q) (\phi(K)^T V)$$

Random Features for Large-Scale Kernel Machines - NeurIPS 2007

- <https://people.eecs.berkeley.edu/~brecht/papers/07.rah.rec.nips.pdf>

Random Feature Attention - ICLR 2021

- <https://arxiv.org/abs/2103.02143>
- <https://arxiv.org/abs/2009.14794>

# Low-Rank Approximations

- Nyström approximation ( - also a kernel method) <https://arxiv.org/abs/2102.03902>

The diagram illustrates the Nyström approximation of a matrix  $B$ . On the left, a large matrix  $B$  (yellow) is partitioned into four blocks:  $K_m$  (orange, top-left),  $S^T$  (blue, top-right),  $S$  (blue, bottom-left), and  $B$  (yellow, bottom-right). This matrix is shown to be approximately equal to the product of three matrices: a tall blue matrix  $S$ , a small orange square matrix  $K_m^\dagger$ , and a wide matrix consisting of an orange block  $K_m$  and a blue block  $S^T$ .

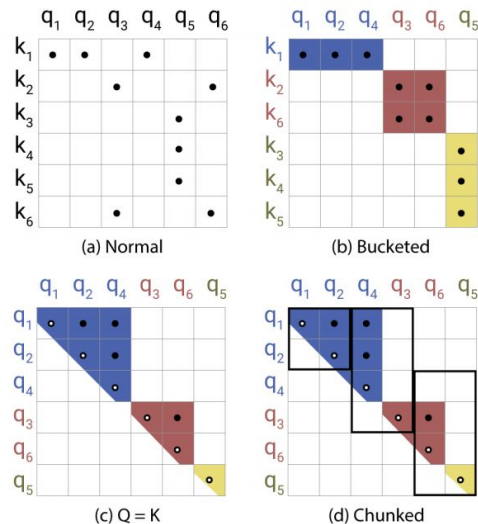
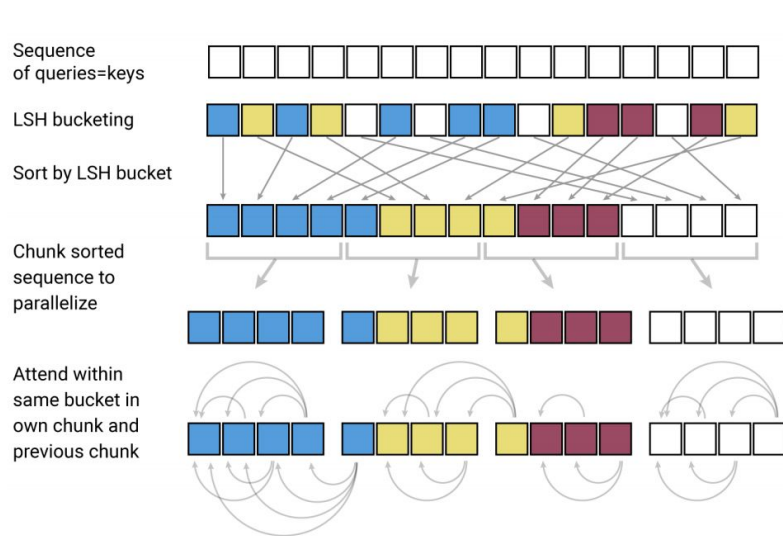
[http://homepage.divms.uiowa.edu/~zli79/talks/dual\\_perspective\\_nystrom\\_method.pdf](http://homepage.divms.uiowa.edu/~zli79/talks/dual_perspective_nystrom_method.pdf)

- JL Lemma: <https://arxiv.org/abs/2006.04768>
  - dimensionality reduction technique
  - random orthogonal projections of a set on points on smaller subspaces preserves L2 norms
  - project valued and keys

$$V \in \mathbb{R}^{N \times d} \rightarrow V' \in \mathbb{R}^{k \times d}$$

$$K \in \mathbb{R}^{k \times d} \rightarrow K' \in \mathbb{R}^{k \times d}$$

# LSH <https://arxiv.org/abs/2001.04451>



## Idea:


1. use locality-based hashing as a clustering mechanism to group related tokens together
2. dot product similarity <https://arxiv.org/pdf/1509.02897.pdf>



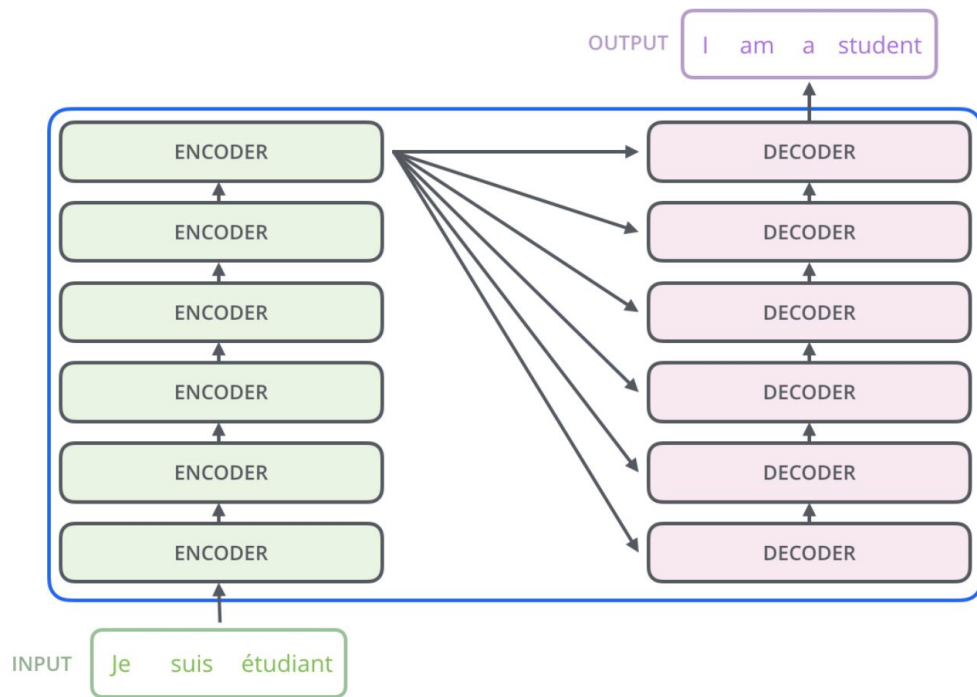
# Transformer Models

# Two Architectures


Encoder Models



Encoder-Decoder Models

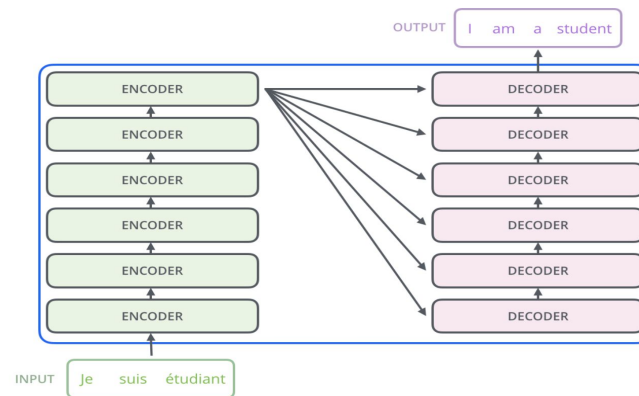


# Three Model Types



**autoregressive**

next token prediction



**seq-2-seq**

autoregressive decoder

**autoencoder**

masked language modeling

# NLP applications

- Language modeling (next token prediction): autoregressive model
- Language understanding: autoencoder model
- Translation: seq2seq

# Method

# Attention Pruning Method

- Train a transformer model
- Perform a forward pass through all the train set
- Gather average attention pattern for each attention mechanism
- Learn which tokens correlate on this problem domain

# Results

# Language Modelling: prune 90% with good performance

$p$ (%)	Perplexity
0	24.157
20	24.157
40	24.214
60	24.566
80	25.115
90	26.011

Transformer-XL-base trained on WikiText-103




# NLU: prune 60% with good performance


Accuracy (%)										
<i>p</i> (%)	MNLI-m	MNLI-mm	QNLI	QQP	RTE	SST-2	MRPC	CoLA	STS-B	Average
0	84.07	83.44	90.91	87.51	65.7	91.97	88.77	57.78	88.39	82.06
20	84.00	83.42	89.72	86.37	64.44	91.11	87.12	55.99	87.34	81.06
40	83.42	83.70	88.76	84.73	62.82	89.91	84.45	52.33	86.48	79.62
50	83.32	82.87	87.81	83.84	62.09	89.05	83.08	48.56	85.28	78.43
60	82.54	81.98	87.19	83.10	61.37	88.82	82.04	45.05	81.70	77.09
80	79.29	78.64	82.37	81.32	57.22	84.52	78.57	34.80	65.89	71.40
90	75.40	75.23	77.23	77.45	49.46	80.56	79.41	20.28	51.39	65.16

BERT base fine tuned on GLUE


# Translation Attention Patterns



encoder-decoder



encoder self-attention



encoder entmax self-attention

# Translation Performance: Cross-Attention is more brittle

$p$ (%)	Self-Enc	Self-Dec	Cross
0		34.94	
20	34.53	34.94	33.50
40	33.70	34.94	24.38
50	33.56	35.08	22.60
60	33.68	34.91	15.08
80	33.67	34.90	6.39

IWSLT14 de-en

Translation: prune\* 60-80% with good performance

$p$ (%)	IWSLT14 de-en	WMT17 en-de
0	34.94	26.73
20	34.92	26.21
40	33.70	26.60
50	33.68	26.19
60	33.64	26.44
80	33.81	21.88

\* prune only self-attention mechanisms

# SQuAD Results <https://www.deepspeed.ai/>

Kernel	Pruning	Efficiency	
	$p$ (%)	Time (s)	Memory (GB)
CUDA	0	95.80	6.24
Triton	0	95.41	6.85
Triton	90	86.44 (↓9.4%)	5.00 (↓27%)

## Performance:

- $p = 0$  yields 81.02 Exact and 88.63 F1
- $p = 90$  yields 79.62 Exact and 87.32 F1

Thank you!